



# Car space occupancy

Niccolò Salvi, Alessio Villa, Beatrice Zani



# Why?

- **Crucial for autonomous vehicles:** collision avoidance, path planning, parking assistance, traffic analysis.
- Need for robust, real-time vehicle **space estimation** where conventional methods fail (**nighttime**, tunnels, poorly lit areas).
- **Final goal:** use main extracted features and geometric properties to build a bounding box around the vehicle.



# Our work

1. Data collection
2. Camera calibration
3. Feature extraction
4. Standard localization via homography
5. Nighttime localization from pair of images
6. Poor perspective localization using out-of-plane features
7. PnP-based vehicle pose estimation from key points


# Data Collection

- **Recording device:** iPhone 13 (main wide-angle lens).
- **Camera setup:** static tripod, 2.5m high, 10m from road.
- **Video settings:** 4K resolution, 30 FPS.
- **Locked auto-exposure (AE) & auto-focus (AF):** to ensure consistent camera properties for accurate geometry

# Camera calibration & CAD

**Data collection & K:** recording **99 diverse frames** of a **standard checkerboard pattern** and processing them with **OpenCV in Python** to determine the camera's unique **intrinsic parameters (K matrix)** and distortion coefficients.




$$K = \begin{bmatrix} 3316 & 0 & 1912 \\ 0 & 3320 & 1072 \\ 0 & 0 & 1 \end{bmatrix}$$



$$\begin{bmatrix} 2.43773846e - 01 \\ -1.59544680e + 00 \\ -1.15284213e - 03 \\ 4.19886247e - 04 \\ 3.56681588e + 00 \end{bmatrix}$$



# CAD Model

**Reference vehicle model:** we found a detailed 3D model of a Škoda Fabia. We exploited this model and extracted exact physical dimensions to better estimate and visualize the 3D bounding box around the car.




# Frames extraction

Counted frames present in our video and extracted a fixed number of frames.



# Taillights extraction


- **Initial Vehicle Localization:** we exploited YOLO to find the vehicle and crop the image, allowing us to focus specifically on the vehicle's rear and reduce background noise
- **Taillight isolation:** convert the image from RGB to **HSV color space** and defined dual hue ranges (near  $0^\circ$  and  $180^\circ$ ) to capture the full spectrum of red light emitted by taillights.
- Morphological operations and contour filtering to clean up the isolated red regions, ensuring we identify the **two largest and most prominent red areas** as the taillights.







# Licence plate extraction

- Defined a **ROI** around the previously located lights.
- We masked the ROI based on **license plate colors** (e.g., white plate body, blue strip)
- Extracted contours from the largest found region.



# Side mirror extraction

- **Vehicle cropping:** expand the YOLO vehicle cropped image to ensure the full side mirror is included.
- Apply **Canny edge detection**.
- **Mirror localization:** filter contours, keeping only the most prominent shapes. The side mirror is then identified by locating the **contour point with the highest (X-axis) position**



# 1. Localization via homography

- Select a single image where both taillights and license plate corners are visible
- Identify four key features:
  - TL and TR = top corners of the license plate
  - L2 and R2 = taillights points
- Compute the direction of rays:
  - Use camera intrinsics  $\mathbf{K}$  to back-project TL and TR to 3D rays.
- Estimate the 3D position of TL and TR by:
  - Measuring angle between rays
  - Applying known real-world plate width constraint

- Estimate vehicle orientation from vanishing point:
  - Compute intersection of two parallel edges (e.g., license plate top and taillights)
  - Back-project VP to get 3D direction (vehicle's forward axis)
- Build a local coordinate system:
  - $X$ : from TL  $\rightarrow$  TR
  - $Y$ : from vanishing direction, orthogonalized to  $x$
  - $Z = X \times Y$
- Position the vehicle origin at plate center, offset vertically
- Generate 3D bounding box using known vehicle dimensions




## 2. Nighttime localization from pair of images

- Select two non consecutive frames.
- Track the two taillights (L1,R1 and L2,R2) and calculate two critical vanishing points:
  - $V_x$ : from the intersection of taillights segments within each frame (horizontal motion).
  - $V_y$ : from the intersection of corresponding taillights across the two frames (depth-wise motion).
- Find **vanishing line** by crossing the two vanishing points found.




- Find the the **normal vector to the vehicle's rear plane** using the vanishing line. This is achieved by multiplying  $\ell$  with the transpose of the camera's intrinsic matrix.
- Compute back-projected rays of the two taillights using  $K^{-1}$ . The **angle** between these two rays is found using geometry rules.
- Use the angle and the known width of the lights to find **distance** of rear plane from lights.
- Exploit known CAD dimension **to draw the bounding box**.



### 3. Poor perspective localization using out-of-plane features

- Aims at solving the localization problem when perspective is lacking.
- reduce the 3D-to-2D projection to a **2D-to-1D problem** by projecting model points onto a **horizontal plane through the camera center**, simplifying the unknown to a single rotation angle ( $\theta$ ).
- Use an **iterative method** to solve the problem.




# Iterative method

- **Initial estimate:** start with pose of two rear lights and two symmetric corners of the licence plate.
- **Yaw angle:** yaw angle ( $\theta$ ) is found by computing **vertical vanishing point** and **horizontal vanishing line**. Symmetric image points are projected onto the vanishing line to precisely calculate  $\theta$ .
- **Loop:**
  - update vehicle lateral direction using the current  $\theta$
  - perform triangulation using updated rays and known distances
  - refine the back-plane normal via cross product or SVD
  - rotate the back-plane normal to compute the new vertical direction
  - recompute vanishing geometry and find new estimate of  $\theta$
- Iteration 1:  $\theta = 6.27^\circ$ ,  $\Delta\theta = -2.34^\circ$
- Iteration 2:  $\theta = 5.89^\circ$ ,  $\Delta\theta = -1.88^\circ$
- Iteration 3:  $\theta = 5.58^\circ$ ,  $\Delta\theta = -1.55^\circ$
- Iteration 4:  $\theta = 4.37^\circ$ ,  $\Delta\theta = -1.28^\circ$
- Iteration 5:  $\theta = 4.17^\circ$ ,  $\Delta\theta = -0.21^\circ$
- Iteration 6:  $\theta = 4.14^\circ$ ,  $\Delta\theta = -0.04^\circ$



# Result




## 4. PnP-based vehicle pose estimation from key points

- **PnP:** Perspective-n-Point (PnP) aims at estimating the **3D position and orientation** of a camera relative to a known 3D object.
- Achieved using two rear lights points and two licence plate points.
- advantage of PnP over homography methods is its ability to handle **non-coplanar 3D points**.
- **Variants:**
  - **Iterative PnP:** refines an initial guess for higher accuracy.
  - **EPnP (Efficient PnP):** provides a computationally efficient and stable non-iterative solution, ideal for real-time.
  - **PnPRANSAC:** combines PnP with outlier detection, robustly handling noisy or incorrect point matches.


# 4 points approach

We evaluated both an iterative PnP solver (with an initial pose estimate) and a direct PnP solution (like EPnP).

EPnP




Iterative



# 5 points approach

We incorporated a fifth point to find the estimate, to make it more robust. The side mirror incorporates crucial **depth variation**, which enables more accurate estimation of the car's lateral extent.



# Comparison of results

Methods 1 and 4 are based on basic geometric transformations and image matching techniques, making them computationally efficient and easy to implement. However, they face significant limitations in terms of generalizability and reliability.




## Method 1

- Simple and fast to implement
- Highly dependent on planar assumptions
- Performs poorly under perspective distortion or non-flat surfaces

## Method 4

- Easy to implement
- Fast runtime
- Very limited robustness to pose and scale variation
- Not reliable under real-world conditions



# Comparison of results

Methods 2 and 3 leverage more advanced geometric reasoning and perspective cues to estimate the orientation of objects with higher robustness.




## Method 2

- High accuracy with well-detected vanishing points
- Robust to perspective changes
- Depends on clear line features and image resolution

## Method 3

- Handles diverse and noisy scenes
- Learns robust features, works well across diverse conditions
- Computationally demanding



# Future work

- Automate vehicle data retrieval using a neural network for license plate recognition.
- Query public databases to identify car model.
- Access 3D CAD repositories for accurate geometry.
- Enhance robustness of plate and light detection under challenging conditions with advanced object detection algorithms.